

Marketing Cloud Readiness Overview

Driving Readiness, Maximizing Success in the Cloud.



Marketing Cloud Lifecycle Development: Sandbox & Multi-Environment Guide

Audience: Internal Cloud Success Resources

Purpose: Help teams guide customers through sandbox strategy, lifecycle testing, and multi-environment promotion for Marketing Cloud Next (MCN) campaigns, Data Cloud, and Agentforce.

Part 1: Understanding the Sandbox Landscape

Platform / Core Sandbox Types

Marketing Cloud Next runs natively on Core (Salesforce Platform), which means the standard Core sandbox model applies. The type of sandbox determines what a marketer can test and what data they can work with.

Sandbox Type	What It Copies	Best For Marketing Use
Developer Sandbox	Production metadata (config) only — lower storage limits	Initial flow/segment build & testing with synthetic data
Developer Pro Sandbox	Production metadata only — larger storage than Dev	Integration testing, segment logic, UAT with moderate data
Partial Copy Sandbox	Full metadata + sample of production data (defined by template)	UAT, segment testing against representative audience data
Full Copy Sandbox	Entire production org — all metadata + all data	Stress testing, complex multi-channel flows, load testing,

Key principle: Marketing Cloud Next does not have a traditional sandbox model like MCE (where every business unit is effectively "production"). MCN leverages the Core sandbox infrastructure, which means you get true isolated environments for lifecycle development.

What Does a Marketing Sandbox Give You?

In a Core sandbox with MCN provisioned:

- Full Campaign, Flow, Segment, and Content Builder capability
- Data Cloud provisioned automatically (metadata replicated, **data is NOT copied**)
- Email Blackhole capability (Spring '26 GA — see Part 4)
- Ability to test flows, segments, IR rules, and advanced actions without production risk

Part 2: Data Cloud in the Lower Environment

What a Data Cloud Sandbox IS

A Data Cloud sandbox is an isolated tenant provisioned from the same production org. When a sandbox is created or refreshed, Data Cloud automatically provisions with:

-  All **configuration/metadata** replicated (Data Streams, DMOs, Calculated Insights, Segments, Transformations, Identity Resolution rulesets)
-  Access to Data Cloud features: Segmentation, Activation, Calculated Insights, Data Graphs, Vector DB, Zero Copy (BYOL)
-  Mergeback support via Change Sets, Metadata API, and DevOps Center

What a Data Cloud Sandbox IS NOT

-  **Data is NOT copied from production.** You must re-authenticate connectors and re-ingest data.
-  **Connections/authentications are NOT replicated.** While connector metadata (configuration) does copy over to the sandbox, any connector that is external to the Core environment — such as Snowflake, S3, or Marketing Cloud Engagement — ***requires an admin to manually enter new authentication credentials to activate it in the lower environment.*** Without this step, the connector metadata exists but no live connection is established, meaning viable lower environment testing against those external systems is not possible until re-authentication is complete.
-  **Not free.** Sandbox activity draws from the same shared credit pool as production (see below).

After refreshing, complete Data Cloud provisioning by clicking **Get Started** in Data Cloud Setup, then re-

authenticate or add new connections to hydrate the sandbox with data.

Data Cloud Sandbox Credit Consumption

As of **June 17, 2025**, there is no longer a separate Data Cloud Sandbox SKU. Key pricing realities:

- **Sandbox and production credits are merged into a single shared pool.** No separate sandbox purchase needed.
- **Sandbox consumption rate = 20% lower than production rates** — Salesforce applies a lower multiplier to encourage dev/testing activity.
- **Plan for 10–20% of your production credit estimate for sandbox use** (typical guideline for representative testing).
- **Sandbox storage now counts toward combined total** — customers should be aware this may affect storage reporting.
- Monitor production vs. sandbox consumption separately via **Digital Wallet**.

"In Sandbox environments, consumption during the activation phase is significantly lower than in production. For instance, if production scoping estimates 30M credits for segmentation or data queries, Sandbox usage for the same would likely be no more than 10% of that, as the activity isn't happening at production scale." [1](#)

Data Cloud Rate Card (Production vs. Sandbox)

Usage Type	Unit	Production Credits	Sandbox Credits (20% lower)
External Data Pipeline (ingest)	Per 1M rows processed	2,000	~1,600
Data Transforms	Per 1M rows processed	400	~320
Data Query (standard)	Per 1M rows processed	2	~1.6
Segmentation / Activation	Per 1M rows processed	Varies	~20% reduction
Streaming (External)	Per 1M rows	5,000	~4,000
Zero Copy / Federation	Per 1M rows accessed	70	~56

Consumption credits are purchased in bundles of 100,000 credits for \$500. Sandbox refresh restores metadata but **not data** — all data must be reloaded and reprocessed, which can create a significant credit

impact. Plan refreshes accordingly.

Part 3: Email & Messaging Sending Costs in Lower Environments

Marketing Cloud Next — Email Sends

Environment	Send Behavior	Credit Impact
Production	Real sends to real recipients	Full email/messaging send credits consumed
Sandbox (standard)	Real sends to configured recipients	Send credits ARE consumed Email Blackhole is enabled
Sandbox + Email Blackhole	Messages suppressed unless domain is on allow list	No email send credits consumed for suppressed messages

Email Blackhole (Spring '26 GA)

Email Blackhole is a guardrail feature available in MCG/MCA sandboxes that prevents accidental sends to real customers during testing, training, or UAT.

- Admin configures up to **5 allow-listed domains** — only messages to those domains are delivered
- All other messages are **suppressed and logged** with "not sent" status
- **Suppressed messages do NOT consume email credits** — protecting sender reputation and send allotment
- Great for: volume stress testing, flow execution testing, complex multi-channel campaign builds

Important nuance: For legacy MCE blackhole sending, sends **do count toward super message allotment** — this is different from the MCG/MCA blackhole behavior.

Messaging Credits in Lower Environments

- **SMS, WhatsApp, Push** sends in sandbox follow the same consumption logic as email — they will consume messaging credits unless suppressed.
- Use Email Blackhole (where applicable) or test with internal allow-listed domains to minimize credit impact.
- For stress testing complex multi-channel campaigns, combine Email Blackhole with representative (but not production-scale) data sets.

Part 4: What Can (and Cannot) Be Promoted Across Environments

This is the most critical concept for marketers to understand in a multi-environment lifecycle.

What CAN Be Promoted (Metadata / Infrastructure)

These are **structural components** — they define *how* a campaign runs, not the data it runs against. They are deployable via Change Sets, Data Kits, DevOps Center, or Metadata API.

Component	Promotion Method	Notes
Campaign Flows (segment-triggered, event-triggered, on-demand)	Change Set / DevOps Center, Metadata API and CLI	Flow logic, wait steps, branch conditions
Segments (configuration/definition)	Data Kit using Change Set	Segment rules and logic — Not underlying audience data
Calculated Insights	Data Kit using Change Set	
Email Content / Templates	Content is in the DigitalExperience component and is promoted using Change Set / DevOps Center, Metadata API and CLI	<i>template structure</i> can be promoted
Identity Resolution Rulesets	Data Kit using Change Set	IR rules, reconciliation logic
Data Streams (configuration)	Data Kit using Change Set	Connection config is not included — must be re-authenticated
DMOs / Data Model	Data Kit using Change Set	Mappings, transformations, calculated insights
Data Graphs	Data Kit using Change Set	Graph configurations
Agentforce Agents / Topics / Actions	Change Set / DevOps Center	Agent metadata, prompt templates, instructions
Flows (Standard Salesforce)	Change Set / DevOps Center, Metadata API and CLI	Subflows, suppression flows, advanced action flows

Permission Sets	Change Set / DevOps Center, Metadata API and CLI	User access configurations
Activation Targets	Data Kit (roadmap for some connectors)	Some connectors require manual setup



Recommended Promotion Sequence — Deploy in This Order

Tip from the field: The order you deploy matters — skip a step and the deployment will fail, not because your work is wrong, but because the environment isn't ready to receive it yet.

Think of it like setting a dinner table. You need the tablecloth down before the plates, and the plates down before the glasses on top. Each layer depends on what came before it.

Always deploy in this order:

1. **Data Cloud objects** (*the foundation*) DMOs, Data Streams, Calculated Insights, Identity Resolution rulesets, and Data Graphs. Everything else references these — they have to exist in the target environment first, full stop.
2. **Content** (*the materials*) Email templates, dynamic content blocks, and DigitalExperience components. Your Flow is going to call on these to send messages — if the content isn't there yet when the Flow arrives, it has nothing to point to.
3. **Flow** (*the wiring*) Campaign Flows, subflows, suppression flows, and advanced action flows. By the time the Flow lands, all of its dependencies — the data objects it reads from and the content it sends — are already in place and waiting.

The plain-language reason this matters: MCN validates dependencies at deployment time. If a Flow references a Data Cloud object or a content template that doesn't exist yet in the target environment, the deployment fails immediately. It's not a bug — it's the platform checking its work. Deploy in dependency order, and it goes smoothly every time.

...

✗ What CANNOT Be Promoted (Data)

These are **data records** — they exist as specific instances tied to a given environment. They cannot be moved between orgs.

Component	Why It Cannot Be Promoted
Specific email send records	Transactional data — environment-specific

Segment population / audience lists	Data (not config) — must be re-built by running the segment in the target org
Data Cloud ingested data	Data records live in the tenant — must be re-ingested in each environment
Contact engagement history	Engagement records are data, not metadata
Send logs / reporting data	Reporting metrics are environment-specific
Identity Resolution output (unified profiles)	Computed data — re-runs in each environment based on ingested data
Calculated Insight results	Computed data — re-runs in each environment

The Core Mental Model: Think of it this way — the *blueprint* (flow logic, segment rules, template structure) moves between environments. The *results of running that blueprint* (who is in the segment today, what emails were sent, unified profile outputs) are data and stay in each environment.

Part 5: Email Campaign Lifecycle — Multi-Environment Journey

This section walks through how an email campaign travels through a realistic multi-environment development lifecycle.

Example: Promotional Email Campaign with IR Rule Testing

Business Goal: Test a new IR (Intelligent Recommendation) rule that provides different value propositions to audience populations based on behavioral data, then validate the flow logic before production deployment.

Environment Architecture

DEV Sandbox → QA/UAT Sandbox → Staging (Full Copy) → Production

Stage 1: DEV Sandbox — Build

What you do here:

- Create the **segment definition** (rules based on behavioral/engagement DMOs)

- Build the **campaign flow** structure: wait steps, branch logic, send email actions
- Create **email content templates** (layout, placeholder personalization tokens)
- Configure the **IR rule** logic in the flow (path split based on profile attributes)
- Set up **test data connections** (re-authenticate Data Cloud connectors with sandbox data sources)

What you DON'T do here:

- Send to real audiences
- Rely on production-scale identity resolution output
- Use full production data volumes

Cost profile: Minimal — limited data ingestion, no/low send volume



Stage 2: QA / UAT Sandbox — Test IR Rule Logic

This is the ideal environment to stress-test **which populations receive which value proposition** based on your IR rule.

What you test:

- Does the IR rule correctly branch contacts into different paths?
- Does the segment filter produce the expected population subsets?
- Do wait steps and exit rules behave as designed?
- Do personalization tokens resolve correctly for each branch?

Email Blackhole: Enable this so you can **execute the flow at volume** without sending real emails or consuming send credits. Flow execution logic, branching, timing, and IR rule evaluation all run normally — only the actual send is suppressed.

What you promote to get here:

- Flow (via Change Set / DevOps Center)
- Segment definition (via Data Kit)
- Email templates (via Change Set)
- IR rule configuration (via flow/metadata promotion)

What you rebuild in this environment:

- Data connections (re-authenticate)
- Seed the org with representative test data (re-ingest)
- Re-publish segments to generate audience population



Stage 3: Full Copy Sandbox (Staging) — Stress Test Complex Campaigns

The Full Copy sandbox is where **complex multi-channel campaigns and advanced architecture** should be validated before production.

When to use this tier:

- Multi-channel campaign (email + SMS + WhatsApp + agent handoff)

- Complex flow with subflows, suppression logic, advanced wait conditions
- Agentforce agent integration (e.g., forward to agent action after email engagement)
- High-volume stress testing with production-like data
- Load testing segment execution and flow concurrency

What you validate:

- End-to-end flow execution at volume
- Agent handoff logic — does the "Forward to Agent" action trigger correctly for the right segment?
- Multi-step flows with interleaved events (email open → wait → SMS fallback)
- Suppression subflow execution at scale
- Data Graph performance at production-representative volumes

Email Blackhole: Highly recommended here. You can execute the **full message-sending architecture** — including multi-channel — without consuming production send credits or messaging credits. Use your internal domain allow list to spot-check actual send rendering.

Cost profile: Higher than DEV/QA due to larger data volumes and more compute — plan for ~15–25% of production credit estimates for a thorough staging run.

Stage 4: Production

What you promote from Staging:

- Validated flow (via DevOps Center / Change Set)
- Final email templates (via Change Set)
- Segment definitions (via Data Kit — re-publishes in prod against live audience)
- IR rule configuration
- Agent configuration (if applicable)

What you build fresh in production:

- Activate campaign
- Publish segment to generate live population
- Go live with real send execution

What Makes This Different from Traditional MCE

In legacy Marketing Cloud Engagement (MCE), there is effectively no concept of UAT vs. Production — every business unit is "production." This means:

- **MCN on Core is a fundamental shift** — marketers now have true lifecycle management
 - Infrastructure (flows, segments, content templates) can travel across environments
 - Data (audience records, send history, engagement data) stays in each environment
 - This creates a real dev→test→staging→prod cycle for the first time in Marketing Cloud
-

Part 6: Advanced Testing Patterns

Pattern 1: IR Rule A/B Testing in Lower Environment

Use case: Test whether a new IR rule produces meaningfully different engagement predictions across audience segments before production deployment.

Setup:

1. Build two flow variants (or use Path Experiment element) — one with existing rule, one with new IR rule
2. Seed sandbox with representative behavioral data
3. Run both branches with Email Blackhole enabled
4. Analyze which segment population falls into which branch
5. Validate segment sizes and path logic before deploying to production

Promotes to production: Flow with finalized IR rule branch logic, segment definition

Pattern 2: Complex Multi-Channel Stress Test with Email Blackhole

Use case: Before launching a complex campaign (email + SMS + agent handoff), validate infrastructure at volume.

Setup:

1. Build entire campaign architecture in Full Copy Sandbox:
 - a. Email flow with multiple content variations
 - b. SMS fallback for non-openers (event-triggered sub-flow)
 - c. Agentforce agent handoff action for high-engagement responders
 - d. Suppression subflow to prevent messaging opted-out contacts
2. Enable Email Blackhole — configure internal domain(s) as allow list
3. Load representative volume of test contacts (10K–100K+ depending on target production volume)
4. Execute the full flow — all messaging suppressed except to internal domains
5. Review:
 - a. Flow execution logs
 - b. Branch outcomes by population
 - c. Agent handoff trigger rate
 - d. Suppression subflow execution
 - e. Exit rule evaluation

Why this is highly valuable: You can test the entire architecture — including agent handoffs and multi-channel messaging logic — without consuming production email/SMS credits or risking customer messaging.

Promotes to production: Entire flow architecture, email templates, SMS configurations, agent setup (all as metadata)

Pattern 3: Segmentation Testing Across Environments

Use case: Validate that a segment definition produces the expected population before activating in production.

Key principle: The **segment rules** (metadata) travel — the **population results** (data) are computed fresh in each environment.

Steps:

1. Build segment in DEV with synthetic data — validate rule logic
 2. Promote segment definition (Data Kit) to QA
 3. Hydrate QA with sample data (Partial Copy or re-ingested representative data)
 4. Publish segment in QA — compare population size/composition against expectations
 5. Promote to production — publish in prod against live unified profiles
-

Part 7: Key Principles for Cloud Success Teams

Top Guidance Points for Customer Conversations

1. **Sandboxes are now a real thing for marketing customers** — MCG/MCA on Core means true lifecycle management exists for the first time in Marketing Cloud.
 2. **Data Cloud sandbox ≠ free.** Credits draw from the shared pool at 20% lower rates. Plan for 10–20% of production estimates for sandbox credit usage.
 3. **Email Blackhole is a game-changer for stress testing.** Customers can now run full flow execution at volume — including multi-channel and agent handoffs — without consuming send credits. This is Spring '26 GA.
 4. **What travels across environments: flow logic, segment rules, content templates, IR configurations (metadata).** What stays in each environment: audience data, engagement history, send records (data).
 5. **Data Cloud sandbox requires manual re-hydration.** Metadata copies automatically; connections and data do not. Budget time and credits for re-authentication and re-ingestion.
 6. **Full Copy Sandbox is the right tier for complex Agentforce + marketing flow testing.** Agents are non-deterministic — you need production-like data volumes and infrastructure to validate behavior.
 7. **Sandbox refresh = credit impact.** Every sandbox refresh triggers re-ingestion. Plan refresh cadence carefully and avoid unnecessary refreshes during active development.
-

Part 8: The MCE Marketer's Translation Guide — From "Fake" Lower Environments to Real Lifecycle Management

This section is written specifically for **MCE customers and their teams** who are navigating a fundamental shift in how environments work. If your team has lived in Marketing Cloud Engagement and used Package Manager as a promotion tool, this is your orientation to what changes, what's better, and what requires a genuinely different mindset.

The MCE Reality: Production All the Way Down

In Marketing Cloud Engagement, there was no true concept of a lower environment. Every Business Unit (BU) operated as a fully capable, fully live, sending-capable environment. Some customers created dedicated BUs — sometimes even entire secondary Enterprise tenants — and used the **MCE Package Manager** to shuttle assets between them. This worked. Teams built their workflows around it. But it was, at its core, a workaround dressed up as a lifecycle.

Here's what made it feel like a lower environment:

- A separate BU (or secondary tenant) would be designated as "Dev" or "UAT"
- Package Manager could export and import **Journeys, Email content, Data Extensions (including their data), Automations, and supporting assets** — essentially most of the building blocks of a campaign
- Teams would build and test in this BU, verify behavior, then promote the package to the production BU

Here's what made it *not* a real lower environment:

- **There was no off switch for sending.** Every BU was production-capable. If a Journey was accidentally activated in the "Dev" BU, it would send. To real people. With real consequences.
- **Suppression lists were BU-specific.** Opt-outs and suppression data didn't automatically carry across — meaning your "lower" BU could message someone who had unsubscribed in production.
- **Data Extensions with data** could be promoted — which meant teams were sometimes working with copies of real customer data in ostensibly non-production environments, with the governance and compliance risks that come with that.
- **There were no compute or access boundaries** — no concept of storage limits, org isolation, or environment-level permissions enforced by the platform. You were managing all of this manually.

In short: MCE's "lower environments" were production environments that teams had agreed, by convention, not to use for production purposes. The platform itself had no knowledge of the distinction.

What CRM Sandboxes Actually Are (And Why It's Different)

When MCN marketers encounter Core sandboxes for the first time, the most important reframe is this:

CRM sandboxes are not "extra orgs you don't send from." They are platform-enforced isolated environments with real architectural boundaries, intentional limitations, and defined promotion gates.

This is a fundamentally different contract between the marketer and the platform.

Dimension	MCE "Lower Environment" (Package Manager BU)	CRM Core Sandbox
Sending capability	Always on — full production capable	Sending active, but Email Blast can suppress all non-allowlist sends
Data isolation	Manual — you choose what to copy	Platform-enforced — data do NOT copy on sandbox create/refresh
Metadata/config promotion	Package Manager exports most assets including data	Change Sets, Data Kits, Dev Center — metadata only, no c
Data Extension / audience data	Could be included in Package export	Cannot be promoted — must ingested per environment
Suppression/compliance	BU-specific — gaps between environments possible	Org-level — sandbox is truly isolated from production con
Environment type hierarchy	Flat — all BUs equal capability	Tiered — Dev, Dev Pro, Partial Copy, Full Copy carry different capabilities and copy behavior
Storage/compute limits	Shared across account — no environment-level limits	Sandbox-type-specific limits enforced by the platform
Refresh behavior	N/A — BUs don't "refresh"; they persist indefinitely	Sandboxes can be refreshed from production, restoring metadata to known state
Governance controls	Manually enforced by team convention	Platform-enforced isolation, permission sets, access control environment

What Package Manager Could Promote vs. What MCN Promotion Supports

This is where the most friction will surface for MCE-experienced marketers. Package Manager was remarkably broad in what it could move — and that breadth created a sense that "everything travels." MCN's promotion

model is more intentional: **only the blueprint travels; the execution context stays behind.**

Asset Type	MCE Package Manager	MCN Promotion (Change Set Data Kit)
Email content / templates	✅ Full email content, HTML, dynamic content blocks	✅ Template structure (DigitalExperience component content promotes)
Journey / Flow logic	✅ Journey Builder flows with steps and configuration	✅ Campaign Flows via Change Set / DevOps Center
Data Extensions (schema)	✅ DE structure (field definitions)	✅ DMO / Data Model schema Data Kit
Data Extension data (records)	✅ Could include data rows in export	❌ Data does NOT promote — must be re-ingested in target environment
Automations	✅ Automation Studio schedules and activities	✅ Standard Salesforce Flow (equivalent automation) via Change Set
Segments	⚠️ Filtered DEs (audience lists as data)	✅ Segment <i>definition</i> (rules) Data Kit — population re-run in target env
Suppression lists	⚠️ Could be promoted as a DE with data	❌ Not directly portable — compliance data is environment-specific
Connector / integration config	⚠️ Some via installed packages	✅ Connector metadata via Data Kit — authentication must be entered
Send logs / engagement data	❌ Not promotable	❌ Not promotable — environment-specific data
User permissions / roles	⚠️ Limited, mostly manual	✅ Permission Sets via Change Set

The critical mindset shift: In MCE Package Manager, a marketer could export a Journey *and* the Data Extension *with its subscriber data* and have a working, populated test environment almost immediately. In MCN, the Journey (flow) and the segment definition (rules) promote — but the audience data must be re-ingested and the segment must re-publish in the new environment. The *structure* arrives; the *population* must be built.

Part 9: 📁 Environment Purpose & Content Recommendations — MCE Marketer's Lens

For marketers coming from MCE, the question isn't just "what can I do here?" — it's "what *should* I be doing here, and does this match how fast I need to move?" The table below maps each sandbox type to its ideal use, what content to build or load there, and what kind of testing is genuinely viable.

Environment	Marketer's Mental Model	What to Build / Load Here	Viable Testing	Not Appropriate For
Developer Sandbox	"My drafting table" — lowest fidelity, fastest to set up	Flow logic skeletons, segment rule definitions, email template structure, synthetic test data (hand-crafted or generated)	Logic validation: do branches fire correctly? Do wait steps resolve? Does personalization syntax work?	Volume testing, suppression validation, data dependent segment population testing
Developer Pro Sandbox	"My integration workbench" — more room, same metadata-only starting point	All of Dev, plus connector re-authentication, moderate synthetic dataset, integration configurations	Connector behavior, Data Cloud stream validation, segmentation logic against structured (non-production) data	Compliance testing, off, production representative volume, stakeholder testing
Partial Copy Sandbox	"My UAT environment" — closest MCE-familiar feel, real data shape without full exposure	Promoted flow + segment metadata, re-authenticated connectors, Partial Copy data set (production-sampled by template)	Segment population against representative audience, stakeholder review, content QA with realistic data, compliance team review	Full-scale load testing, agent behavior validation at volume

Full Copy Sandbox	"My dress rehearsal" — this is what MCE never had	Full promoted architecture, production-replica data, all connectors re- authenticated	End-to-end volume testing, multi- channel execution with Email Blackhole, agent handoff validation, stress testing, final compliance sign- off before go-live	None — this highest-fidel pre-producti environment available
-------------------	---	--	---	--

Governance, Compliance & the Immediacy Tension

One of the real tensions for MCE marketers moving to MCN is the gap between **marketer speed** and **lifecycle discipline**. MCE made it easy to move fast — build in one BU, package it, deploy it, send. The review gates were largely self-imposed.

In MCN on Core, the platform enforces the gates. This is good for compliance-heavy organizations. It's an adjustment for fast-moving marketing teams.

Low governance / high immediacy teams may want to consider:

- A lean two-environment model: a single Dev/UAT sandbox + Production, with Email Blackhole as the send guardrail
- Streamlined Change Set deployments for well-understood campaign patterns (e.g., recurring promotional sends)
- Dev Sandbox as the primary build environment, Partial Copy for stakeholder sign-off only

High governance / compliance-driven teams (regulated industries, large enterprise, data-sensitive verticals) should leverage the full model:

- Four-tier lifecycle: Dev → Dev Pro or Partial Copy → Full Copy Staging → Production
- Formal promotion gates at each tier with documented approval
- Partial Copy or Full Copy sandboxes used for compliance team review before any production deployment
- Separation of concerns: marketing team builds in Dev/Dev Pro; compliance and legal review in Partial Copy; final validation in Full Copy
- Suppression list validation at the Full Copy tier before production promotion

The core governance reframe for MCE teams: In MCE, governance was something your *team* imposed on the platform. In MCN on Core, governance is something the *platform* supports your team in enforcing. The sandbox tiers, the promotion gates, the isolation boundaries — these are infrastructure for responsible marketing execution, not bureaucratic overhead. The discipline pays off in compliance confidence, reduced production incident risk, and a clear audit trail of what was validated before it went live.

Quick Reference: "I Used to Do This in MCE – What Do I Do Now?"

MCE Pattern	MCN Equivalent	Key Difference
Build in Dev BU → Package Manager export → Import to Prod BU	Build in Dev Sandbox → Change Set / Data Kit promotion → Production org	Data does not travel with the promotion — must re-ingest in environment
Use a secondary Enterprise tenant as a "UAT environment"	Partial Copy or Full Copy Sandbox provisioned from production org	Platform-enforced isolation; limits by sandbox type; connections require re-auth
Test a Journey by activating it in a Dev BU and watching sends	Test a Flow in sandbox with Email Blackhole enabled — full execution, sends suppressed	Blackhole suppresses sends by default; no risk of accidentally emailing real customers
Copy a Data Extension with subscriber data to lower BU for testing	Re-ingest a representative dataset into the sandbox Data Cloud tenant	Data is never copied automatically — re-authentication and re-ingestion is always required
Promote a filtered DE as your "segment" to another BU	Promote segment <i>definition</i> (rules) via Data Kit; re-publish in target environment	The audience list is computed in each environment — only the rules travel
Manage suppression lists per BU	Suppression is org-level in MCN; sandbox is fully isolated from production contacts	No risk of sandbox activity affecting production suppression state

Part 10: Known Limitations & In-Progress Items

Maintenance note: This section is reviewed and updated on a per-release basis. If you encounter behavior that conflicts with what's documented here, check the release notes for the most recent version or reach out to the MC Readiness team.

Not everything described in this guide works end-to-end today. This section documents confirmed gaps between the *intended* promotion experience and what's *currently functional* — so you can set accurate expectations with customers and know what's on the roadmap.

 Invocable Actions: Promotion Limitations

The **Send to MCE Email** and **Send to Journey** invocable actions are intended to be fully promotable via Change Sets. Today, that's partially true — but there are meaningful restrictions depending on how the action is configured.

Action	Current Promotion Status	Details
Send to MCE Email (empty/template action)	✅ Promotable	Empty actions with no configuration. Data Extension deployment succeeds today.
Send to MCE Email (configured with a Data Extension)	⚠️ Conditionally promotable	Active deployment fails if the referenced Data Extension has never been used in the target org. Workaround: Deploy as Draft, open in Flow Builder to materialize the DE before activating. Bug record: W-22201680 (P2).
Send to Journey (any configured action)	❌ Not currently promotable	Cannot be promoted from sandbox when configured against a Journey/Data Extension not present in the target org. No current workaround for active deployments. Fix in progress — targeted for release 266.

Root cause (for technical audiences): When a Flow containing a configured invocable Send to MCE Email or Send to Journey action is deployed to a target org where the referenced Data Extension has never been used, the change set deployment fails with a field integrity exception. The platform cannot materialize the Data Extension Business Process Object (BPO) without a prior reference in the target org. Engineering fix estimated at ~10–13 days of work.

 Other Confirmed Promotion & Environment Limitations

Area	Current Status	Details
MCE ↔ Data Cloud connection in lower environments	❌ Cannot connect in sandbox	Data 360 must be provisioned in Production first. A sandbox refresh is required to propagate Data

		tenant availability. No workarounds – this is a Data 360 infrastructure constraint.
Campaign records	✗ Not promotable	Campaign records are data, not metadata. They cannot be deployed via Change Sets, Dev Center, or Copado and must be manually recreated in each target environment.
MCG/A Foundations Edition in sandbox	✗ Does not work in sandbox	Foundations Edition is not supported in sandbox environments. Growth or Advanced Edition is required for sandbox viability.
MCG/A sandbox provisioning	⚠ Known provisioning bug	Sandbox provisioning of MCG/A can fail in some configurations. Known issue: W-21292105. Consult with the MC Readiness team for current status and workarounds.

✅ What This Means for Your Lifecycle Workflow

These limitations are real, but they don't block the overall value of the lifecycle approach described in this guide. Here's how to work around them today:

- **For Send to MCE Email:** Build and test in sandbox with empty (template) actions. Attach the real Data Extension only after deployment by opening the Flow in the target org's Flow Builder and materializing the DE before activating.
- **For Send to Journey:** Document the flow architecture and journey configuration thoroughly in sandbox. Plan to manually recreate the configured action in production until release 266 delivers the fix. This is actively in progress per the MCE/Flow convergence team.
- **For Campaign records:** Include campaign recreation as a deliberate step in your production deployment runbook – budget time accordingly.
- **For Data 360/MCE sandbox connection:** Ensure the initial MCE ↔ Data Cloud connection is established in Production before provisioning or refreshing sandboxes. Sandbox refresh will then propagate the connection availability.

📚 Resources & References

Resource	Description
Agentforce Development Lifecycle Deep Dive	Full sandbox + deployment lifecycle deck (internal)
Data Cloud Sandbox Refresh & Deployment Strategy	Step-by-step sandbox refresh and deployment guide
Spring '26 Release Overview – Marketing Cloud	Email Blackhole, Flow Enhancements, MCE+ details
Data Cloud Best Practices: Consumption Reduction	Consumption optimization guidance for sandbox
Marketing Cloud Consumption & Digital Wallet	Credit tracking, custom tagging, consumption monitoring